

# penalized\_logistic\_regression.R

Erica Zhou

Nov 16, 2022

```
Evaluate <- function(true_label, pred_label) {
  # Compute the 0-1 loss between two vectors
  #
  # @param true_label: A vector of true labels with length n
  # @param pred_label: A vector of predicted labels with length n
  # @return: fraction of points get misclassified

  #####
  # TODO #
  #####

  error <- sum(true_label != pred_label)/length(true_label)

  #####
  #                END OF YOUR CODE #
  #####
  return(error)
}

Predict_logis <- function(data_feature, beta, beta0, type) {
  # Predict by the logistic classifier.
  #
  # Note: n is the number of examples
  #       p is the number of features per example
  #
  # @param data_feature: A matrix with dimension n x p, where each row corresponds to
  #   one data point.
  # @param beta: A vector of coefficients with length equal to p.
  # @param beta0: the intercept.
  # @param type: a string value within {"logit", "prob", "class"}.
  # @return: A vector with length equal to n, consisting of
  #   predicted logits,      if type = "logit";
  #   predicted probabilities, if type = "prob";
  #   predicted labels,      if type = "class".

  n <- nrow(data_feature)
  pred_vec <- rep(0, n)

  #####
  # TODO #
  #####
```

```
#####

if (type == "logit"){
  pred_vec <- as.vector(beta%*%data_feature + beta0)
}

if (type == "prob"){
  pred_vec <- as.vector((exp(
    beta0 + data_feature%*%beta))/(1 + exp(
    beta0 + data_feature%*%beta)))
}

if (type == "class"){
  pred_vec <- as.numeric(((exp(
    (beta0 + data_feature%*%beta))) / (1 + exp(
    (beta0 + data_feature%*%beta))) >= .5))
}

#sample.data  sample.data <- matrix(c(1,3,4,1, 5,8,9,6, 1,3,3,8), 4,3)

#####
#                               END OF YOUR CODE                               #
#####

return(pred_vec)
}

Comp_gradient <- function(data_feature, data_label, beta, beta0, lbd) {
  # Compute and return the gradient of the c
  #
  # Note: n is the number of examples
  #       p is the number of features per example
  #
  # @param data_feature: A matrix with dimension n x p, where each row corresponds to
  #   one data point.
  # @param data_label: A vector of labels with length equal to n.
  # @param beta: A vector of coefficients with length equal to p.
  # @param beta0: the intercept.
  # @param lbd: the regularization parameter
  #
  # @return: a (p+1) x 1 vector of gradients, the first coordinate is the gradient
  #   w.r.t. the intercept.

  n <- nrow(data_feature)
  p <- ncol(data_feature)
  grad <- rep(0, 1 + p)

  #####
  # TODO:                                     #
  #####

```

```

grad0 <- exp(beta0)/(1 + exp(beta0))/n
grad1 <- 1/n*((-data_label + t(exp(beta0 + data_feature**beta)/(1 + exp(
  beta0 + data_feature**beta)))) **data_feature + lbd**beta)
grad <- c(grad0, grad1)

#####
#                               END OF YOUR CODE                               #
#####
return(grad)
}

Comp_loss <- function(data_feature, data_label, beta, beta0, lbd) {
  # Compute and return the loss of the penalized logistic regression
  #
  # Note: n is the number of examples
  #       p is the number of features per example
  #
  # @param data_feature: A matrix with dimension n x p, where each row corresponds to
  #   one data point.
  # @param data_label: A vector of labels with length equal to n.
  # @param beta: A vector of coefficients with length equal to p.
  # @param beta0: the intercept.
  # @param lbd: the regularization parameter
  #
  # @return: a value of the loss function

  #####
  # TODO:                                     #
  #####
  n <- length(data_label)
  p <- exp(beta0 + data_feature**beta)/(1 + exp(beta0 + data_feature**beta))

  loss <- as.numeric((data_label**log(p)+(1-data_label)**log(1-p))/(-n)
    + lbd/2*(norm(beta,"2"))^2)

  #####
  #                               END OF YOUR CODE                               #
  #####
  return(loss)
}

Penalized_Logistic_Reg <- function(x_train, y_train, lbd, stepsize, max_iter) {
  # This is the main function to fit the Penalized Logistic Regression
  #
  # Note: n is the number of examples
  #       p is the number of features per example
  #
  # @param x_train: A matrix with dimension n x p, where each row corresponds to
  #   one training point.

```

```

# @param y_train: A vector of labels with length equal to n.
# @param lbd: the regularization parameter.
# @param stepsize: the learning rate.
# @param max_iter: a positive integer specifying the maximal number of
#   iterations.
#
# @return: a list containing four components:
#   loss: a vector of loss values at each iteration
#   error: a vector of 0-1 errors at each iteration
#   beta: the estimated p coefficient vectors
#   beta0: the estimated intercept.

p <- ncol(x_train)

# Initialize parameters to 0
beta_cur <- rep(0, p)
beta0_cur <- 0

# Create the vectors for recording values of loss and 0-1 error during
# the training procedure
loss_vec <- rep(0, max_iter)
error_vec <- rep(0, max_iter)

#####
# TODO:                                     #
# Modify this section to perform gradient descent and to compute #
# losses and 0-1 errors at each iterations.                       #
#####

for (i in 1:max_iter){

  beta_cur <- beta_cur - stepsize * Comp_gradient(x_train, y_train,
                                                  beta_cur, beta0_cur,
                                                  lbd)[-1]
  beta0_cur <- beta0_cur - stepsize * Comp_gradient(x_train, y_train,
                                                  beta_cur, beta0_cur,
                                                  lbd)[1]

  # 0-1 error: the true vs. the predicted
  # First we need the predicted label for each beta and beta0 starting from (0,0)
  y_pred <- Predict_logis(x_train,beta_cur,beta0_cur, "class" )
  # Second we calculate and add the 0-1 error for each iteration to the vector
  error_vec[i] <- Evaluate(y_train, y_pred)
  # loss of each iteration
  loss_vec[i] <- Comp_loss(x_train, y_train, beta_cur, beta0_cur, lbd)
}

#####
#                                     END OF YOUR CODE                                     #
#####

```

```
return(list("loss" = loss_vec, "error" = error_vec,  
          "beta" = beta_cur, "beta0" = beta0_cur))  
}
```